

OpenCL Basics

Parallel Computing on GPU and CPU

23. März 2011 | Willi Homberg

Agenda

- Introduction
- OpenCL architecture
 - Platform model
 - Execution model
 - Memory model
 - Programming model
- Platform layer
- Runtime
- Language, compiler
- Example: Vector addition

Introduction


- OpenCL – Open Computing Language
 - Open, royalty-free standard
 - For cross-platform, parallel programming of modern processors
 - An Apple initiative
 - Approved by NVIDIA, Intel, AMD, IBM, ...
 - Specified by the Khronos group
- Intended for accessing heterogeneous computational resources
 - CPUs (Intel, AMD, IBM Cell BE, ...)
 - GPUs (Nvidia GTX & Tesla/Fermi, AMD/ATI 58xx, ...)
 - Future processors with integrated graphics chip (Sandy Bridge, Llano)
- Difference to CUDA or CAL/IL
 - Code hardware agnostic, portable

OpenCL Information

- Khronos OpenCL - <http://www.khronos.org/opencvl>
 - Specification and headers files
 - Online manual pages, quick reference card
- MacResearch OpenCL tutorials - <http://www.macresearch.org/opencvl>
- NVIDIA devel. zone: OpenCL - <http://developer.nvidia.com/object/opencvl.html>
 - Best practices guide:
http://www.nvidia.com/content/cudazone/CUDABrowser/downloads/papers/NVIDIA_OpenCL_BestPracticesGuide.pdf
- AMD/ATI – ATI Stream Computing:
 - <http://ati.amd.com/technology/streamcomputing/opencvl.html>
 - SDK - <http://developer.amd.com/gpu/ATIStreamSDK>
- IBM OpenCL for Linux on Power - <http://www.alphaworks.ibm.com/tech/opencvl>


OpenCL Architecture: Platform Model

- A host is connected to one or more (possibly heterogeneous) OpenCL devices
 - A device is divided into compute units (CUs)
 - CUs are subdivided into processing elements (PEs)



OpenCL Architecture: Platform Model

- An OpenCL application runs on the host
 - Submits commands to execute on the PEs within a device
 - PEs execute a single stream of instructions as
 - SIMD units
 - In lockstep with a single stream of instructions
 - SPMD units
 - PEs maintain own program counter



JuGiPSy – Juelich's GPU System

- NVIDIA Tesla 1070
 - 4 x Tesla C1060 GPU
 - 4 x 4 GB DDR3 memory
 - 4 x 102 GB/s memory bandwidth
 - 4 TFLOPS single-precision
 - 346 GFLOPS double-precision
 - 700W
- Host System (2x)
 - Xeon E5430@2,66 GHz quad-core
 - 32 GB memory, 32 KB data cache
 - 1 TB disk

Bandwidth:	MB/s
Host to Device:	2157.0
Device to Host:	1886.4
Device to Device:	73458.1

2x PCIe 16x



Transtec 1000R S1070

JUGIPSY

- Network
 - External interface: jugipsy.zam.kfa-juelich.de
 - Internal private 1 GE network:
 - gipsy1, gipsy2
 - each connected to 2 GPUs
- Software
 - OpenSuSE 11.1
 - NVIDIA
 - NVIDIA_GPU_Computing_SDK (CUDA-3.1)
 - AMD / ATI
 - ati-stream-sdk-v2.2-lnx64



JUDGE – Juelich Dedicated GPU Environment

Heterogeneous cluster: 60 TFLOPS peak

54 IBM System x iDataPlex dx360 m3

each:

- 2 NVIDIA Tesla M2050 (Fermi)
 - 1,15 GHz (448 cores), 3 GB memory
 - 148 GB/s memory bandwidth
 - 1,03 TFLOPS single-precision
 - 515 GFLOPS double-precision
 - 225W
- 2 Intel Xeon X5650(Westmere)
 - 6-core processor 2,66 GHz
 - 96 GB memory



Bandwidth:	MB/s
Host to Device:	3422.5
Device to Host:	2911.1
Device to Device:	85636.2

JUGIPSY: OpenCL Devices

PLATFORM NAME : ATI Stream PLATFORM VERSION : OpenCL 1.0 ATI-Stream-v2.1 (145)

PLATFORM PROFILE : FULL_PROFILE

PLATFORM VENDOR : Advanced Micro Devices, Inc.

1 devices found supporting OpenCL

Device name : Intel(R) Xeon(R) CPU E5430 @ 2.66GHz

Max Compute Units : 4

Amount of Global Memory : 3221225472 bytes

Amount of Local Memory : 32768 bytes

Max Work Group Size : 1024

Max Work Item Dimensions : 3

Max Work Items on dimension 0 : 1024

Max Work Items on dimension 1 : 1024

Max Work Items on dimension 2 : 1024

=====

PLATFORM NAME : NVIDIA CUDA PLATFORM VERSION : OpenCL 1.0 CUDA 3.0.1

PLATFORM PROFILE : FULL_PROFILE

PLATFORM VENDOR : NVIDIA Corporation

2 devices found supporting OpenCL

Max Compute Units : 30

Amount of Global Memory : 4294770688 bytes

Amount of Local Memory : 16384 bytes

Max Work Group Size : 512

Max Work Item Dimensions : 3

Max Work Items on dimension 0 : 512

Max Work Items on dimension 1 : 512

Max Work Items on dimension 2 : 64

JUDGE: OpenCL Devices

Device 0: "Tesla M2050"

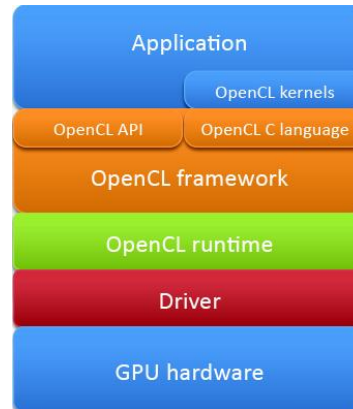
CUDA Driver Version:	3.20
CUDA Runtime Version:	3.20
CUDA Capability Major/Minor version number:	2.0
Total amount of global memory:	2817982464 bytes
Multiprocessors x Cores/MP = Cores:	14 (MP) x 32 (Cores/MP) = 448 (Cores)
Total amount of constant memory:	65536 bytes
Total amount of shared memory per block:	49152 bytes
Total number of registers available per block:	32768
Warp size:	32
Maximum number of threads per block:	1024
Maximum sizes of each dimension of a block:	1024 x 1024 x 64
Maximum sizes of each dimension of a grid:	65535 x 65535 x 1
Maximum memory pitch:	2147483647 bytes
Texture alignment:	512 bytes
Clock rate:	1.15 GHz
Concurrent copy and execution:	Yes
Run time limit on kernels:	No
Integrated:	No
Support host page-locked memory mapping:	Yes
Compute mode:	Default (multiple host threads can use this device simultaneously)
Concurrent kernel execution:	Yes
Device has ECC support enabled:	Yes
Device is using TCC driver mode:	No

Device 1: "Tesla M2050"

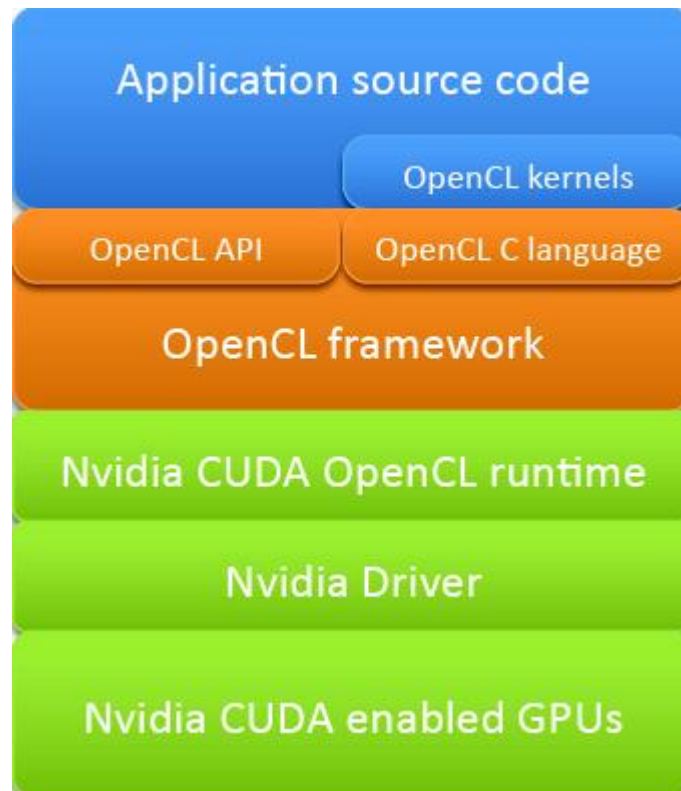
OpenCL Architecture: Execution Model

- Compute Kernel
 - Runs on the OpenCL device
 - Basic unit of executable code (similar to a C function)
 - Data-parallel or task-parallel
- Compute Program
 - Collection of compute kernels and internal functions
 - Analogous to a dynamic library
- Applications
 - Run on the host
 - Queue compute kernel execution instances
 - Instances are executed in-order or out-of-order
 - Events are used to implement appropriate order of execution

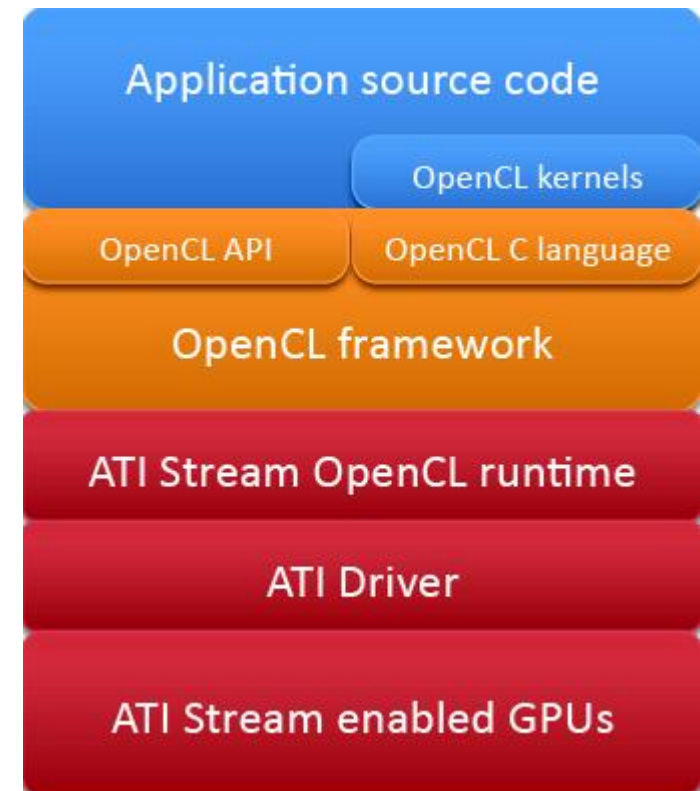
OpenCL Stack (GPUs)



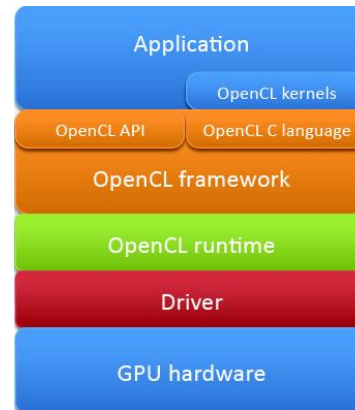
NVIDIA-GPU



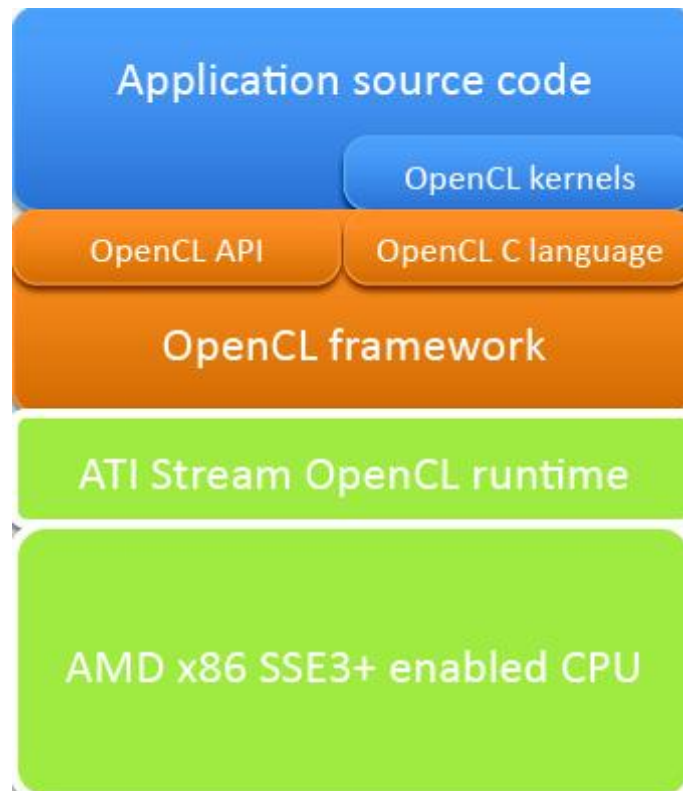
AMD-GPU



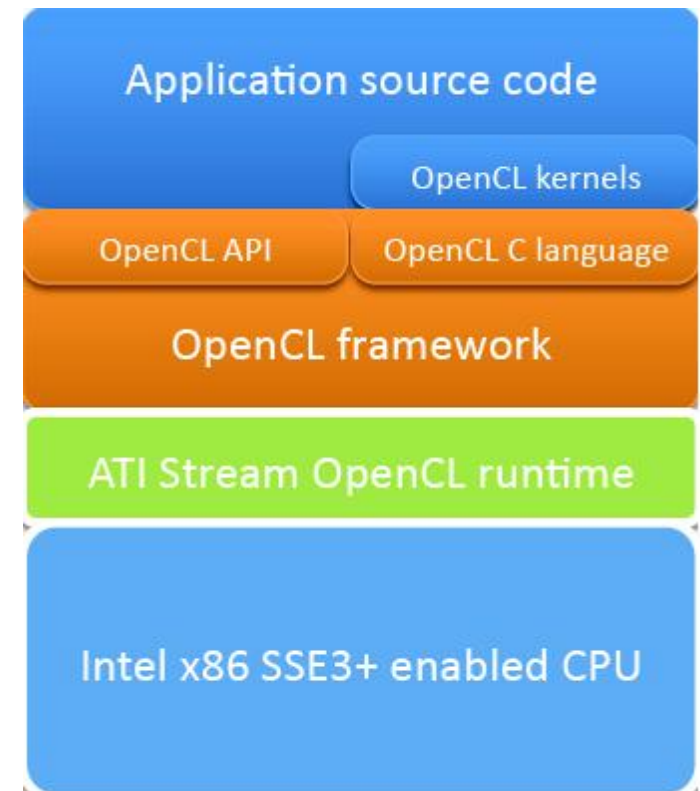
OpenCL Stack (CPUs)



AMD-CPU



Intel-CPU



OpenCL Architecture: Execution Model

Program execution occurs in two parts

- Host program executes on the host
 - Defines contexts for kernels and manages execution
 - Data parallel programming model
 - NDRange index space (1-, 2- or 3-dimensional)
 - Task parallel programming model
- Kernels execute on one or more devices
 - An instance of the kernel is called **work-item** (CUDA: thread)
 - Executes for each point in index space
 - Organized in **work-groups** (CUDA: blocks)

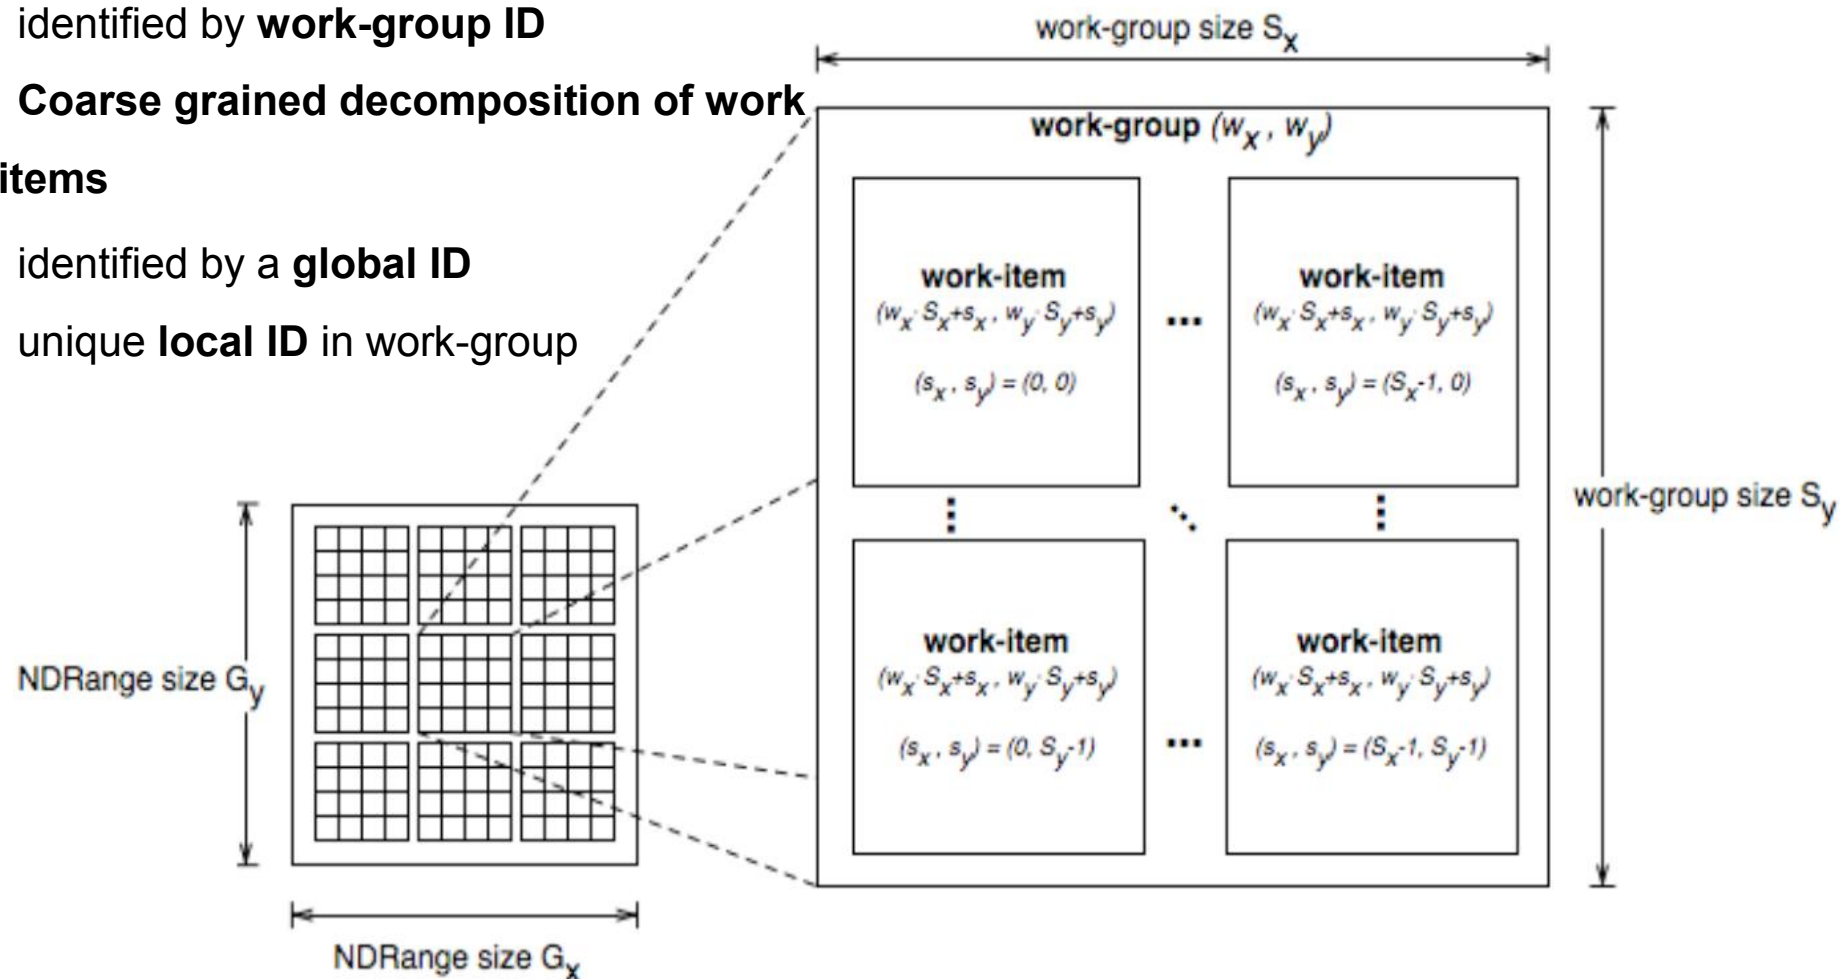
OpenCL Architecture: Execution Model

■ Work-groups

- identified by **work-group ID**
- **Coarse grained decomposition of work**

■ Work-items

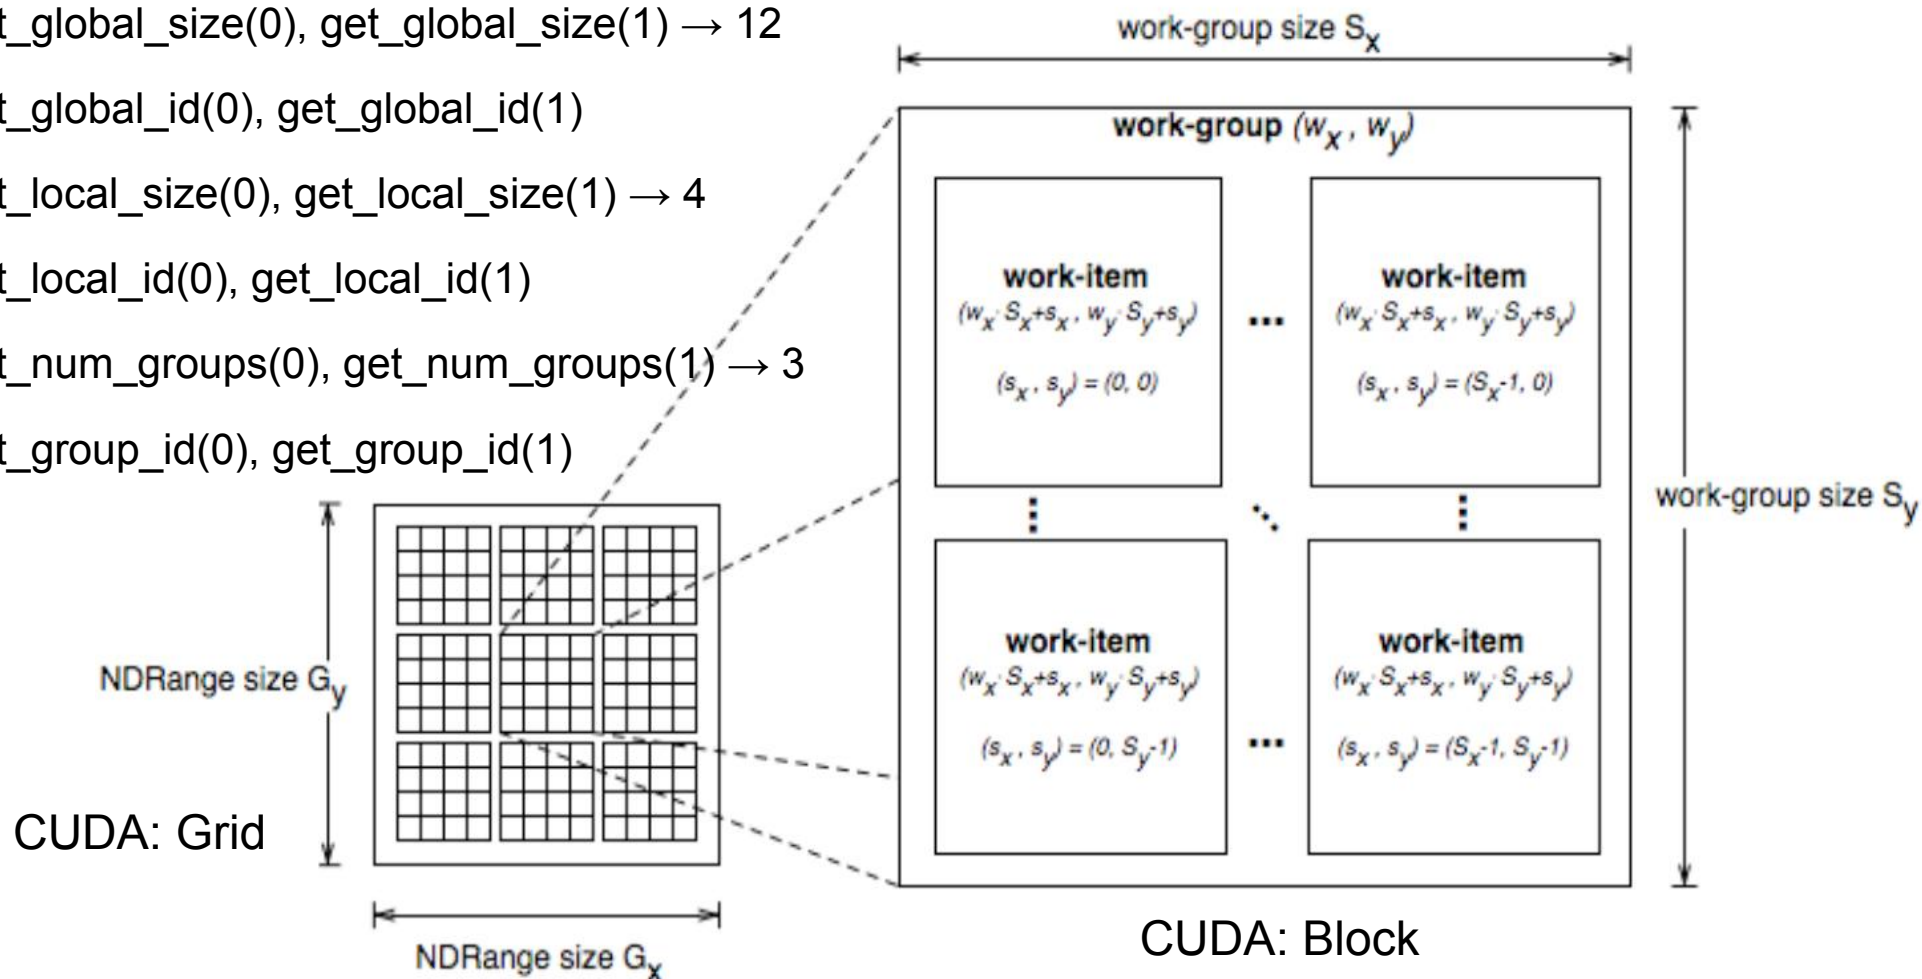
- identified by a **global ID**
- unique **local ID** in work-group



OpenCL Architecture: Execution Model

Built-in functions: `get_work_dim()` $\rightarrow 2$


- `get_global_size(0)`, `get_global_size(1)` $\rightarrow 12$
- `get_global_id(0)`, `get_global_id(1)`
- `get_local_size(0)`, `get_local_size(1)` $\rightarrow 4$
- `get_local_id(0)`, `get_local_id(1)`
- `get_num_groups(0)`, `get_num_groups(1)` $\rightarrow 3$
- `get_group_id(0)`, `get_group_id(1)`



Kernel: Vector Addition

```
__kernel void
vec_add (__global const float *a,
         __global const float *b,
         __global float *c)
{
    int gid = get_global_id(0);

    c[gid] = a[gid] + b[gid];
}
```



OpenCL Architecture: Execution Model

- Host defines context for execution of the kernels
- Context includes
 - Devices
 - Collection of OpenCL devices to be used by the host
 - Kernels
 - OpenCL functions that run on OpenCL devices
 - Program objects
 - Program source and executable that implement the kernels
 - Memory objects
 - Visible to host and devices
 - Values that can be operated on by instances of a kernel

OpenCL Architecture: Execution Model

- Command queues
 - Data structure to coordinate execution of kernels on the devices
 - Commands are placed in queue and scheduled onto the devices
 - Kernel execution command (execute a kernel)
 - Memory commands (transfer data)
 - Synchronisation commands
 - Scheduling properties in-order, out-of-order
 - Event objects control execution
 - Refer to kernel execution or memory commands
 - Coordinate execution between host and devices and between commands
 - Multiple queues with a single context
 - Run concurrently, no synchronisation

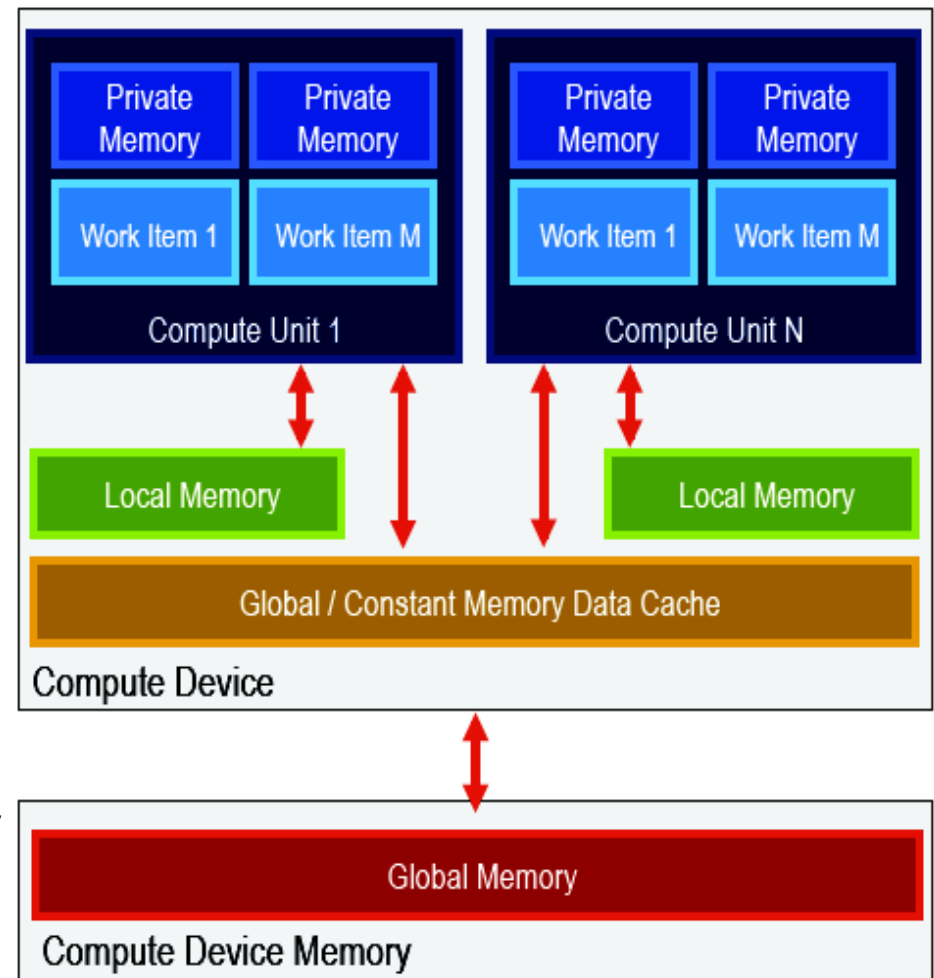
OpenCL Architecture: Memory Model

- Work-items executing a kernel have access to four memory regions
 - Global memory – read/write access to all work-items
 - Constant memory – remains constant during the execution of the kernel
 - Local memory – region local to a work-group (CUDA: shared)
 - Private memory – region private to a work-item (CUDA: local)

	Global	Constant	Local	Private
Host	Dynamic allocation	Dynamic allocation	Dynamic allocation	No allocation
	Read / Write access	Read / Write access	No access	No access
Kernel	No allocation	Static allocation	Static allocation	Static allocation
	Read / Write access	Read-only access	Read / Write access	Read / Write access

OpenCL Architecture: Memory Model

- Multiple distinct address spaces
- Address qualifiers
 - `__private`
 - `__local`
 - `__constant`
 - `__global`
- Relaxed consistency
 - Local memory in a single work-group
 - Global memory at a work-group barrier
- Host creates memory objects in global memory
- Memory transfer blocking or non-blocking
- Host may map a region into its address space



OpenCL Architecture: Programming Model

- Data-parallel execution
 - Sequence of instructions applied to multiple elements of a memory object
 - Index space defines how the data maps onto work-items
 - No one-to-one mapping between work-item and element in memory object
 - Programmer defines total number of work-items to be executed in parallel
 - Explicit model – division into work-groups is specified by programmer
 - Implicit model – division into work-groups is managed by OpenCL
- Task-parallel execution
 - Compute devices such as CPUs can also execute task-parallel compute kernels
 - Executes as a single work-item per core
 - A compute kernel written in OpenCL
 - A native C / C++ function

OpenCL Platform Layer

- Implements platform-specific features allowing applications for
 - Querying OpenCL platforms
 - clGetPlatformIDs: provides list of available platform
 - clGetPlatformInfo: provides specific information about platform
 - Getting device information
 - clGetDeviceIDs: provides list of available devices
 - clGetDeviceInfo: provides specific information about device
 - Creating OpenCL contexts
 - clCreateContext: creates an OpenCL context with one or more devices
(for managing command-queues, memory, kernels)
 - clCreateContextFromType: creates context from a device type
 - clGetContextInfo: provides information about a context

OpenCL Runtime, Buffer Objects, Program Objects

- API calls for executing a kernel, reading or writing a memory object
 - Command Queues
 - `clCreateCommandQueue`: creates queue on a specific device
 - Memory objects
 - `clCreateBuffer`: creates a buffer object
 - `clEnqueueReadBuffer`: enqueue read from a buffer to host memory
 - `clEnqueueWriteBuffer`: enqueue write to a buffer from host memory
 - Program objects
 - `clCreateProgramWithSource`: creates program object for context and loads the source code from text strings
 - `clCreateProgramWithBinary`: creates program object for context and loads binary into the program object

OpenCL Runtime, Program Objects, Kernel&Event Objects

- Program objects (cont'd)

- `clBuildProgram`: builds (compiles & links) a program executable for all devices in context associated with program
- `clCreateKernel`: creates a kernel object
- `clSetKernelArg`: set the value for a specific argument of a kernel

- Executing kernels

- `clEnqueueNDRangeKernel`: enqueues a command to execute a kernel using a given range of work-items
- `clEnqueueTask`: enqueues a command to execute a kernel using a single work-item

- Flush and Finish

- `clFlush`: issues all previously queued commands in queue
- `clFinish`: blocks until all queued commands have completed

OpenCL Programming Language

- Derived from ISO C99
- A few restrictions
 - Recursion, function pointers, functions in C99 standard headers, ...
- Built-in Data Types
 - Scalar and vector data types, structs, pointers, data-type conversion functions, ...
- Built-in Functions (see reference card)
 - work-item functions
 - math.h
 - read and write image
 - relational
 - geometric functions
 - synchronization functions

Example: Vector Addition

Compute $c = a + b$ (a, b, and c are float vectors of length N)

```
void VectorAddHost(const float* a, const float* b, float* c, int N){  
    int i;  
    for (i = 0; i < N; i++)  
    {  
        c[i] = a[i] + b[i];  
    }  
}
```

OpenCL Host program:

- Query compute devices
- Create context
- Create memory objects
- Compile and create kernel
- Issue commands to command-queue
- Synchronization of commands
- Clean up resources

Platform, Device, Context

```
cl_platform_id cpPlatform;           // OpenCL platform
cl_device_id cdDevice;               // OpenCL device
cl_context cxGPUContext;             // OpenCL context
cl_command_queue cqCommandQueue;     // OpenCL command queue
cl_int ciErr1;                      // Error code var

// Get an OpenCL platform
ciErr1 = clGetPlatformIDs(1, &cpPlatform, NULL);
// Get a GPU device
ciErr1 = clGetDeviceIDs(cpPlatform, CL_DEVICE_TYPE_GPU, 1, &cdDevice, NULL);
// Create the context
cxGPUContext = clCreateContext(0, 1, &cdDevice, NULL, NULL, &ciErr1);

// Create a command-queue
cqCommandQueue = clCreateCommandQueue(cxGPUContext, cdDevice, 0, &ciErr1);
```

Command-queue, buffer memory objects

```
int N = 11444777;           // Length of float arrays (odd # for illustration)
void *srcA, *srcB, *dst;    // Host buffers for OpenCL test
size_t szGlobalWorkSize;    // 1D var for Total # of work items
size_t szLocalWorkSize;    // 1D var for # of work items in the work group

// set Global and Local work size dimensions
szLocalWorkSize = 256;
// GlobalWorkSize rounded up to the nearest multiple of the LocalWorkSize
szGlobalWorkSize = ((N/(int)szLocalWorkSize)+1)*(int)szLocalWorkSize;

// Allocate host buffers
srcA = (void *)malloc(sizeof(cl_float) * szGlobalWorkSize);
srcB = (void *)malloc(sizeof(cl_float) * szGlobalWorkSize);
dst = (void *)malloc(sizeof(cl_float) * szGlobalWorkSize);
```

Command-queue, buffer memory objects (cont'd)

```
cl_mem cmDevSrcA;           // OpenCL device source buffer A
cl_mem cmDevSrcB;           // OpenCL device source buffer B
cl_mem cmDevDst;            // OpenCL device destination buffer
```

```
// Allocate buffer memory objects
```

```
cmDevSrcA = clCreateBuffer(cxGPUContext, CL_MEM_READ_ONLY,
                           sizeof(cl_float) * szGlobalWorkSize, NULL, &ciErr1);
cmDevSrcB = clCreateBuffer(cxGPUContext, CL_MEM_READ_ONLY,
                           sizeof(cl_float) * szGlobalWorkSize, NULL, &ciErr2);
cmDevDst = clCreateBuffer(cxGPUContext, CL_MEM_WRITE_ONLY,
                          sizeof(cl_float) * szGlobalWorkSize, NULL, &ciErr2);
```

Build program and create kernel

```
const char* cSourceFile = "VectorAdd.cl";           // kernel source file
cl_program cpProgram;                               // OpenCL program

// Read the OpenCL kernel in from source file
char** program_src;
std::vector<std::string> prog_lines;
readCLFile (prog_lines, cSourceFile);
program_src = (char**) malloc (sizeof(char*) * prog_lines.size());
for (unsigned int i = 0; i < prog_lines.size(); i++)
    program_src[i] = (char*) prog_lines[i].c_str();

// Create the program
cpProgram = clCreateProgramWithSource(cxGPUContext, prog_lines.size(),
                                     (const char **)program_src, NULL, &ciErr1);

// Build the program (compile & link)
ciErr1 = clBuildProgram(cpProgram, 0, NULL, NULL, NULL, NULL);
```


Build program and create kernel

```
cl_kernel ckKernel;                                // OpenCL kernel
// Create the kernel
ckKernel = clCreateKernel(cpProgram, "VectorAdd", &ciErr1);

// Set the Argument values
ciErr1 = clSetKernelArg(ckKernel, 0, sizeof(cl_mem), (void*)&cmDevSrcA);
ciErr1 |= clSetKernelArg(ckKernel, 1, sizeof(cl_mem), (void*)&cmDevSrcB);
ciErr1 |= clSetKernelArg(ckKernel, 2, sizeof(cl_mem), (void*)&cmDevDst);
ciErr1 |= clSetKernelArg(ckKernel, 3, sizeof(cl_int), (void*)&N);

// Asynchronous write of data to GPU device
ciErr1 = clEnqueueWriteBuffer(cqCommandQueue, cmDevSrcA, CL_FALSE, 0,
                               sizeof(cl_float) * szGlobalWorkSize, srcA, 0, NULL, NULL);
ciErr1 |= clEnqueueWriteBuffer(cqCommandQueue, cmDevSrcB, CL_FALSE, 0,
                               sizeof(cl_float) * szGlobalWorkSize, srcB, 0, NULL, NULL);
```

Build program and create kernel

```
// Launch kernel (1-dimensional range)
ciErr1 = clEnqueueNDRangeKernel(cqCommandQueue, ckKernel, 1, NULL,
                                &szGlobalWorkSize, &szLocalWorkSize, 0, NULL, NULL);

// Synchronous/blocking read of results, and check accumulated errors
ciErr1 = clEnqueueReadBuffer(cqCommandQueue, cmDevDst, CL_TRUE, 0,
                              sizeof(cl_float) * szGlobalWorkSize, dst, 0, NULL, NULL);

// Cleanup allocated objects
if(ckKernel)clReleaseKernel(ckKernel);
if(cpProgram)clReleaseProgram(cpProgram);
if(cqCommandQueue)clReleaseCommandQueue(cqCommandQueue);
if(cxGPUContext)clReleaseContext(cxGPUContext);
if(cmDevSrcA)clReleaseMemObject(cmDevSrcA);
if(cmDevSrcB)clReleaseMemObject(cmDevSrcB);
if(cmDevDst)clReleaseMemObject(cmDevDst);
```

Kernel: Handle extra indices

```
__kernel void
vec_add (__global const float *a,
         __global const float *b,
         __global float *c,
         int N))
{
    // get index into global data array
    int gid = get_global_id(0);

    // bound check (equivalent to the limit on a 'for' loop for standard/serial C code)
    if (gid >= N)
    {
        return;
    }
    // add the vector elements
    c[gid] = a[gid] + b[gid];
}
```

Array length N	= 11444777
Global Work Size	= 11444992
Local Work Size	= 256
# of Work Groups	= 44707 (last work-group deals only with 215 elements)

Kernel: Use Work-Item Built-in Functions

```

__kernel void
vec_add (__global const float *a,
         __global const float *b,
         __global float *c,
         int N))
{
    // get index into global data array
    // int gid = get_global_id(0);
    int gid = get_local_size(0)*get_group_id(0)
              + get_local_id(0);

    // add the vector elements
    c[gid] = a[gid] + b[gid];
}

```

Array length N	= 8
Global Work Size	= 8
Local Work Size	= 4
# of Work Groups	= 2

e.g. NDRange size Gx=8

float *a=

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

float *b=

7	6	5	4	3	2	1	0
---	---	---	---	---	---	---	---



__kernel void vec_add(...)

float *c=

7	7	7	7	7	7	7	7
---	---	---	---	---	---	---	---

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

7	6	5	4	3	2	1	0
---	---	---	---	---	---	---	---

Work Group 0 Work Group 1
 e.g. work-group size Sx=4

7	7	7	7	7	7	7	7
---	---	---	---	---	---	---	---

JUGIPSY: Getting started

- See login message
 - Functionality tests:
 - /usr/bin/nvidia-smi
 - /srv/NVIDIA_GPU_Computing_SDK/C/bin/linux/release/deviceQuery
 - /srv/NVIDIA_GPU_Computing_SDK/C/bin/linux/release/bandwidthTest
 - For a start, copy the tutorials to your own home directory ...
 - Develop own projects from SDK tutorials
 - NVIDIA:
 - 30 OpenCL tutorials, 71 CUDA tutorials
 - ATI:
 - 37 OpenCL tutorials, 14 CAL tutorials
 - Mind environment, e.g. LD_LIBRARY_PATH sequence:
 - **module load cuda - module switch cuda ati-stream-sdk**

JUDGE: Getting started

Functionality tests:

- `cd /localhome/partec/NVIDIA_GPU_Computing_SDK_3.2.16`
 - `C/bin/linux/release/deviceQuery`
 - `C/bin/linux/release/bandwidthTest`
- For a start, copy the tutorials to your own home directory ...
 - Develop own projects from NVIDIA SDK tutorials (28 OpenCL, 79 CUDA)
- Account application:
 - https://dispatch.fz-juelich.de:8812/account_ident/back=/RESSOURCEN
- Environment: `source NV_env.sh`:

```
export CUDA_HOME=/opt/cuda
export CUDA_INSTALL_PATH=/opt/cuda
export
CUDA_SDK_HOME=/localhome/partec/NVIDIA_GPU_Computing_SDK_3.2.16
export PATH=${CUDA_HOME}/bin:${CUDA_HOME}/computeprof/bin:$PATH
export LD_LIBRARY_PATH=${CUDA_HOME}/lib:$LD_LIBRARY_PATH
export LD_LIBRARY_PATH=${CUDA_SDK_HOME}/lib:$LD_LIBRARY_PATH
export CUDA_VERSION=3.2.16
```

Matrix Multiplication

Matrix sizes:

$wA = 800$

$hA = 1600$

$wB = 800$

Dimensions = 2

Local work sizes: 16x16

Global work sizes: 800x1600


Group sizes: 50x100

`./oclMatrixMul -sizemult=10 -device=0`

`./oclMatrixMul -sizemult=10 -device=0,1`

Device 0

Device 1



Results

```
./oclMatrixMul -sizemult=10 -device=0
./oclMatrixMul Starting... Device 0: Device Tesla T10 Processor
Using Matrix Sizes: A(800 x 1600), B(800 x 800), C(800 x 1600)
```

Running Computations on 1 - 1 GPU's...

MM, Throughput = **175.4612 GFlops/s**, Time = **0.01167 s**, Size = 2048000000, NumDevsUsed = 1, Workgroup = 256
Kernel execution time on GPU 0 : 0.01164 s

Comparing results with CPU computation... PASSED

```
./oclMatrixMul -sizemult=10 -device=0,1
./oclMatrixMul Starting...Device 0: Device Tesla T10 Processor   Device 1: Device Tesla T10 Processor
Using Matrix Sizes: A(800 x 1600), B(800 x 800), C(800 x 1600)
```

Running Computations on 1 - 2 GPU's...

MM, Throughput = **175.4431 GFlops/s**, Time = **0.01167 s**, Size = 2048000000, NumDevsUsed = 1, Workgroup = 256
Kernel execution time on GPU 0 : 0.01166 s

MM, Throughput = **347.5211 GFlops/s**, Time = **0.00589 s**, Size = 2048000000, NumDevsUsed = 2, Workgroup = 256
Kernel execution time on GPU 0 : 0.00587 s
Kernel execution time on GPU 1 : 0.00586 s

Comparing results with CPU computation... PASSED

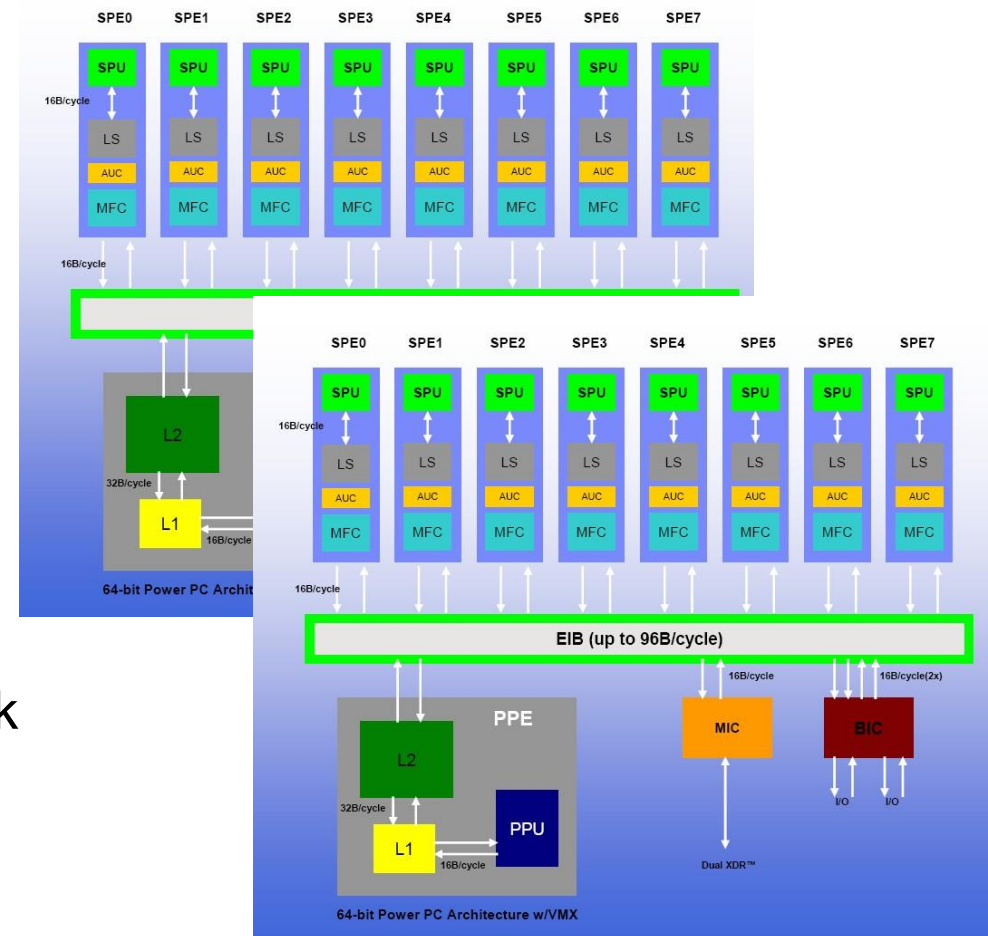
OpenCL Basics

Extra slides

23. März 2011 | Willi Homberg

Cell Cluster JUICEnext

- 35 QS22 blades @ 3.2 Ghz
 - 2 Cell processors (CBE)
 - 8 GB Memory
- 4x InfiniBand adapter (SDR)
- Frontend: x86 compatible
- Interconnect switch: IB (Cisco)
- Peak performance:
 - 1 blade: 217,6 GFLOPS peak (64 bit float)
 - 35 blades: 7 TFLOPS peak (64 bit float)



JUICEnext: Getting started

PLATFORM NAME : IBM
PLATFORM VERSION : OpenCL 1.0 200912040043
PLATFORM PROFILE : FULL_PROFILE
PLATFORM VENDOR : IBM
PLATFORM EXTENSIONS :
2 devices found supporting OpenCL

Device name :ACCELERATOR PowerXCell8i processor
Device version :OpenCL 1.0 200912040043
Device profile :EMBEDDED_PROFILE
Device vendor :IBM
Device extensions :cl_khr_byte_addressable_store
Max Compute Units : 16
Amount of Global Memory : 1939271680 bytes
Amount of Local Memory : 248832 bytes
Max Work Group Size : 256
Max Work Item Dimensions : 3
Max Work Items on dimension 0 : 256
Max Work Items on dimension 1 : 256
Max Work Items on dimension 2 : 256

Device name :ACCELERATOR PowerXCell8i processor
Device version :OpenCL 1.0 200912040043
Device profile :EMBEDDED_PROFILE



JUICEnext: Getting started

OpenCL environment:

access IBM SDK: /srv/Software/OpenCL_Cell
samples: /srv/Software/OpenCL_Cell/samples

e.g. copy sample to your \$HOME directory:

```
cd /srv/Software/OpenCL_Cell/samples  
cp -a blackscholes $HOME
```

then:

```
qsub -l  
cd $HOME/blackscholes/ppc  
make  
./bsop --help
```

